CIS 472/572, Spring 2020
Homework 2 (Programming): Decision Trees
DUE DATE: Submit via Canvas by Saturday, April 25th at 11:55pm.

In this assignment, you will implement the ID3 algorithm for learning decision trees. You may assume that the class label and all attributes are binary (only 2 values). Please use the provided skeleton code in Python to implement the algorithm. You may look at open-source reference implementations, such as WEKA, but please **do not copy code from open-source projects**. Your code must be your own. Undergraduates may complete the assignment in teams of 2. Graduates must complete the assignment alone.

The ID3 algorithm is similar to what we discussed in class: Start with an empty tree and build it recursively. Use information gain to select the attribute to split on. (Do not divide by split information.) Use a threshold on the information gain to determine when to stop. The full algorithm is described in this classic paper (with over 11,000 citations):

http://dept.cs.williams.edu/ andrea/cs374/Articles/Quinlan.pdf

You code should run from the command line on ix and accept the following arguments:

```
./id3 <train> <test> <model>
```

Where train is the name of a file containing training data, test contains test data to be labeled, and model is the filename where you will save the model for the decision tree.

The data files are in CSV format. The first line lists the names of the attributes. The last attribute is the class label.

We are providing skeleton code in Python that handles input, output, and some of the internal data structures. Please use it as the starting point because we'd be using that API to grade.

For saving model files, please use the following format:

```
wesley = 0 :
| honor = 0 :
| | barclay = 0 : 1
| | barclay = 1 : 0
| honor = 1 :
| | tea = 0 : 0
| | tea = 1 : 1
wesley = 1 : 0
```

According to this tree, if wesley = 0 and honor = 0 and barclay = 0, then the class value of the corresponding instance should be 1. In other words, the value appearing before a colon is an attribute value, and the value appearing after a colon is a class value.

Once we compile your code, we should be able to run it from the command line. Your program should take three command line arguments, as shown below:

```
./id3 <training-set> <test-set> <model-file>
```

It should output the accuracy of the decision tree on the test set and write the decision tree in the format described above to the specified file.